



JITing PostgreSQL

Andres Freund

PostgreSQL Developer & Committer

Email: andres@anarazel.de

Email: andres.freund@enterprisedb.com

Twitter: [@AndresFreundTec](https://twitter.com/AndresFreundTec)

anarazel.de/talks/edb-pune-jit-2017-11-03/jit.pdf

Motivation

TPC-H Q01

```
SELECT
  l_returnflag,
  l_linestatus,
  sum(l_quantity) AS sum_qty,
  sum(l_extendedprice) AS sum_base_price,
  sum(l_extendedprice * (1 - I_discount)) AS sum_disc_price,
  sum(l_extendedprice * (1 - I_discount) * (1 + I_tax)) AS sum_charge,
  avg(l_quantity) AS avg_qty,
  avg(l_extendedprice) AS avg_price,
  avg(l_discount) AS avg_disc,
  count(*) AS count_order
FROM lineitem
WHERE l_shipdate <= date '1998-12-01' - interval '74 days'
GROUP BY l_returnflag, l_linestatus
ORDER BY l_returnflag, l_linestatus;
```

What about Parallelism?

Efficiency

Identified Bottlenecks

- plain expression evaluation
- aggregate transition values
- tuple deforming
- grouping comparisons & hashing (aggregation, set functions)
- sorting!

What's JIT?

- create executable code
 - when needed
 - without invoking external program
 - takes time
 - optimizer
- linearize Expression Evaluation
- remove branches from Tuple Deforming
- inline functions
 - no indirect function calls
 - elide function call overhead
 - cross function optimization

Existing Parts

- Expression evaluation machinery (v10)
- Known tupledescs for scan, left, right tree
 - important for JITed deforming
- Jit infrastructure (lifetime, errors, settings, inlining)
 - lifetime: EState integration
 - errors: resowner integration
 - inlining: create bitcode for postgres code, load at JIT time
- Expand executor to use expression engine more
 - Aggregate transition functions
 - grouping comparisons (aggregates, setops, ...)


Old Expression Evaluation Engine

- WHERE a.col < 10 AND a.another = 3
 - ExecEvalExpr()
 - ExecEvalAnd()
 - ExecEvalExpr()
 - ExecEvalFunc() # int4lt
 - ExecEvalScalarVar() # a.col
 - slot_getattr
 - slot_deform_tuple # 1
- ExecEvalExpr()
 - ExecEvalFunc() # int4eq
 - ExecEvalScalarVar() # a.another
 - slot_getattr
 - slot_deform_tuple # 2

New Expression Evaluation Engine

- WHERE a.col < 10 AND a.another = 3
 - EEOP_SCAN_FETCHSOME (deform necessary cols)
 - EEOP_SCAN_VAR (a.col)
 - EEOP_CONST (10)
 - EEOP_FUNCEXPR_STRICT (int4lt)
 - EEOP_BOOL_AND_STEP_FIRST
 - EEOP_SCAN_VAR (a.another)
 - EEOP_CONST (3)
 - EEOP_FUNCEXPR_STRICT (int4eq)
 - EEOP_BOOL_AND_STEP_LAST (AND)

Faster Execution: New Hash-Table & Expression Engine




Tuple Deforming

- Datatypes Known
 - type alignment known
 - possibly previous alignment known
- NOT NULL / NULL is known
- Number of to-be-deformed columns is known
- If there's NOT NULL columns, existence of columns might be known

Faster Execution: JIT Compilation

TPCH Q01 timing

scale 100, fully cached




Example Code

Faster Execution: JIT Compilation

TPC-H Improvements

scale 100, fully cached, no parallelism



Missing Parts

- planner:
 - when to jit and how much to optimize
- executor:
 - OOM handling
 - expression engine for hashing
 - make tupledescs known in more places
 - determine NOT NULL in upper tupledescs
- JIT:
 - correct inlining implementation
 - allow extensions to do inlining
 - improve code generation
 - Caching!